

Harmonyland Project

Bastien Pouëssel, Antoine Boutin, Tom Genlis

Preliminary questions

- 1/ What technical/business constraints should the data storage component of the program architecture meet to fulfill the requirement described by the customer in paragraph «Statistics»? So what kind of component(s) (listed in the lecture) will the architecture need?
- **Persistence capabilities:** Long duration storage
- **Scalability:** Scale horizontally (“daily report will weight 200Gb”)
- The architecture will require **Available** and **Partition tolerant** (AP) Storage
 - **Column oriented databases** (for example Cassandra)
 - Easier add and remove and allow column optimisation/compression

Preliminary questions

- 2/ What business constraint should the architecture need to fulfill the requirement describe in the paragraph «Alert»? Which component to choose?

- Business constraints:

- **Low latency** and **real time alert** to Peacemakers 



Usage of a **stream** (for example Kafka)

(Reporting in the future usage of batch processing)

Preliminary questions

- 3/ What mistake(s) from Harmonystate can explain the failed attempt?
- Harmonystate created a team of **data-scientists only**.
- They didn't hire any **data engineer** who could have be able to set an efficient scalable program
- It has led to difficulties in term of infrastructure and scalability which made the PoC fail

Preliminary questions

- 4/ Harmonystate has likely forgotten some technical information in the report sent by the drone. In the future, this information could help Harmonystate make its harmonywatchers much more efficient. Which information?

- **Timestamp**



allows to make periodic reports and conduct analysis on time series

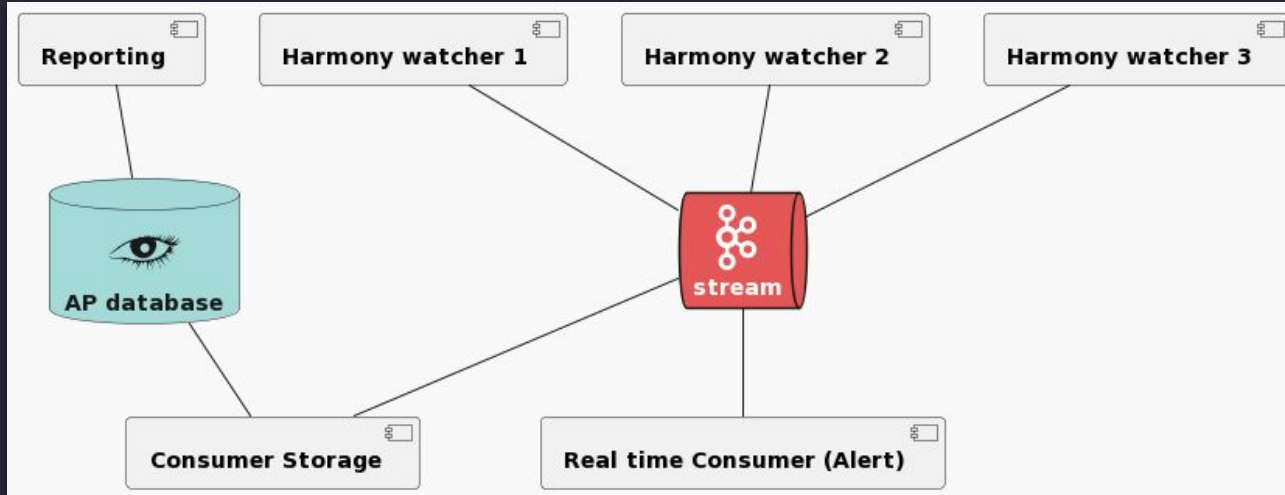


produce any kind of monitoring



failure handling

Architecture





That's all folks !